

Retail Sentiment Intelligence

Real-Time Social Media Mining and Trust-Aware Sentiment Analysis
Using Large Language Models for Retail Feedback Optimization

BITS ZG628T · Dissertation · Post Mid-Semester Presentation

Faculty Mentor: Ms. Pradnya Kashikar

BITS Pilani (WILP)

Supervisor: Mr. Varunendra Pratap Singh

Principal Software Engineer · Walmart Global Tech, Bengaluru

Candidate: Vishal Singh · 2020AA05641

M.Tech (AI & ML) · Birla Institute of Technology & Science, Pilani (WILP)

Agenda

Focus of today — post mid-semester work only

1

Since Mid-Sem — What's New

Post-midsem workstream map

2

Review & Validate (HITL)

Human-in-the-loop correction workflow

3

Smart Reply — Triple Draft

GPT-4o + Mistral 7B + Smart Composer

4

Smart Reply — Prompt & Few-Shot

Prompt template + past validated replies

5

Smart Reply — Worked Example

One complaint → three drafts

6

Learning Loop

Corrections + posted replies → retraining

7

Post Explorer

Multi-facet search across analysed posts

8

Post Lifecycle (Kanban)

Triage → Acknowledged → In-Progress → Resolved

9

Insights & Competitor Analysis

Strategic weekly summaries + cross-brand

10

Notification Centre

In-app + email + Slack, group-routed

11

ModernBERT — Final Results

Macro-F1 0.6272 → 0.7642 (+13.7 pts)

12

Vision — Multi-Pass Payoff

Hallucination 50% → 0%, extraction 25% → 75%

13

Trust-Score Evaluation

15% flagged; 12 of 15 confirmed by annotator

14

Live Demo

Dashboard walkthrough (screenshots)

15

Conclusions & Future Work

RQ1–RQ4 outcomes + realistic roadmap

Since Mid-Sem — What's New

Mid-sem shipped the pipeline; post-midsem made it usable for analysts

DELIVERED AT MID-SEM

- 6-layer offline-first pipeline (25 subreddits)
- ModernBERT fine-tuning up to Stage 2 (F1 0.7285)
- Multi-pass vision with 0% hallucination on pilot
- Trust-score design + weighted formula
- Brand-Health + Alert Feed dashboard pages
- Group-routed email / Slack notifications

ADDED POST-MIDSEM (this deck)

- Review & Validate — human-in-the-loop UI + backend
- Smart Reply Composer — triple-draft (GPT-4o + Mistral + Smart)
- Learning-loop store — corrections & posted replies
- Post Explorer — multi-facet search
- Post Lifecycle — Kanban with 2-step Resolve
- Insights & Competitor Analysis page
- In-app Notification Centre mirroring the digest
- ModernBERT Stage 3 (final) — 0.7285 → 0.7642 (5-fold OOF)
- Trust-score end-to-end evaluation (n = 200)

THE FRAMING QUESTION



Mid-sem answered can the pipeline produce trust-weighted sentiment? Post-midsem answers can an analyst act on it, correct it, and feed the corrections back into the model?

Review & Validate — Human-in-the-Loop

Every prediction can be corrected; every correction is logged

WORKFLOW

- Queue sorted by priority (P1 first) · trust × confidence
- Analyst can override sentiment and aspect tags in one click
- Corrections written to feedback table with analyst + timestamp
- Generate Drafts → three reply options (GPT-4o + Mistral + Smart)
- Analyst picks A, B or C, edits inline, and posts to Reddit
- Posted replies also captured → future few-shot pool

WHY IT MATTERS

Turns a passive dashboard into a training signal — corrections harvested here drive the ModernBERT re-training loop and reply few-shot pool.

NUMBERS TODAY (live capture)

18 posted replies logged · 60 total feedback rows · ~90 % analyst – model agreement on the last 30 reviews.

Retail Sentiment INTELLIGENCE

Brand Health

Lifecycle

Insights

Review

Trust

Alerts

Posts

Pipeline

Notifications

Review

Reddit · Dry-run

Review & Validate

Corrections are saved to the analyses table and feed back into all dashboards.

50 items pending

3 corrections

90.0% agreement

18 replies

SENTIMENT

TIME RANGE

All

Last 7 Days


Apply

MODEL ACCURACY TRACKER

Daily human-agreement rate. Each point = % of reviews on that day where the analyst confirmed the model's sentiment (didn't override it). 30 reviewed to date.

90.0%

OVERALL AGREEMENT



Pending Review

4410

Reviewed

15

neutral

r/OGPBackroom

Confidence: 59%

Trust: 0.55

Correct to:

Positive

Negative

Close

Advanced

It's one more way you can get coached or fired for the same pay as not doing it, I highly recommend just getting used to moving on and accepting the nil picks after you try your best for your own sanity.

4

8/1/2026, 5:49:36 AM

View on Reddit

customer service

online/app

positive

r/samsclub

Confidence: 55%

Trust: 0.46

Correct to:

Neutral

Negative

Close

Advanced

dolphins colors always look good on a chair even if i have no clue who they're playing this season

3

8/1/2026, 5:04:37 AM

View on Reddit

online/app

neutral

r/WalmartEmployees

Confidence: 51%

Trust: 0.54

Correct to:

Positive

Negative

Close

Advanced

Idk

So basically I think i got that parasite. I've been on the toilet for two days, fever of 102 and it's just bad cramps. I had taco bell two weeks ago as well. I have 2.5 points and i'm in the process of being moved over to a different department. If i left to get half a point would that affect it?

1

8/1/2026, 12:10:22 AM

View on Reddit

store experience

online/app

neutral

r/WalmartEmployees

Confidence: 57%

Trust: 0.71

Correct to:

Positive

Negative

Close

Advanced

What does this mean?

I texted my coach about it, this is the same week my approved requested availability for OGP is supposed to start. I gave them the availability of 6am -6pm so I can do school and my other job that evenings. I'm still within my 90 days too. I only have 4 points (a lot of miscommunication and three call out). Someone ease my anxiety. I am full time, before this I was getting 34-38 hours

1

8/1/2026, 12:10:22 AM

View on Reddit

store experience

online/app

4 / 23

Smart Reply Composer — Triple Draft

One prompt, three generators; analyst picks the best draft

Draft A

GPT-4o

Walmart LLM Gateway

Strong reasoning, safest tone
Guardrails via Walmart proxy
~\$0.0002 / reply
Fallback: direct OpenAI

Draft B

Mistral 7B-Instruct

Local Ollama (:11434)

Open-weights, free
Runs fully offline
~15 s warm latency
Strong at retail jargon

Draft C

Smart Composer

Deterministic template

No LLM, always available
Zero latency, zero cost
Content-aware phrase pools
Safety-net fallback

Fallback logic — no draft is ever missing

- GPT-4o unavailable (no gateway key / consumer-ID / network) → Smart Composer draft in slot A
- Mistral unavailable (Ollama not running) → Smart Composer draft in slot B
- Smart Composer always produces slot C — guaranteed reply
- UI labels each card with the actual model used, plus an offline-fallback badge when applicable

Smart Reply — Prompt Design & Few-Shot

Same prompt fed to GPT and Mistral; Smart Composer skips the LLM step

Prompt template (built in `_build_reply_prompt`)

You are a senior Walmart customer-care analyst replying on Reddit.
Write ONE reply to the customer below. Keep it 2-4 sentences,
empathetic, specific to their complaint, no corporate jargon,
no hashtags, no emojis. Do NOT promise refunds you can't verify;
invite them to DM order details if action is needed.
Sign off as a real person, not a brand.

Example customer post: {past_post_1}

Example analyst reply: {past_reply_1}

... (three example pairs total)

Subreddit: r/{subreddit}

Customer ({author}) complaint about: {aspect_1}, {aspect_2}

Customer post: {post_title + post_body, ≤ 1200 chars}

Reply:

PROMPT ANATOMY

① SYSTEM ROLE

who the model is + tone rules

② FEW-SHOT

3 past validated reply pairs
from the feedback table

③ CONTEXT

subreddit · author · top-2 aspects

④ POST

title + body, truncated to 1 200 chars

⑤ TRIGGER

"Reply:" cues the model to generate

Few-shot source — feedback table

- Top-3 past validated replies for the same aspect
- Written by human analysts, already posted to Reddit
- Refreshed on every Generate click → adapts over time

Style guardrails baked into the prompt

- 2 – 4 sentences, empathetic, specific
- No corporate jargon, no hashtags, no emojis
- No unverifiable refund promises → DM to continue
- Sign off as a person, not a brand

Smart Reply — Worked Example (all three drafts real, live-captured)

raw_posts.id = reddit_1u2bgdw · r/samsclub · gateway ✓ · Ollama ✓ · Smart Composer ✓

CUSTOMER COMPLAINT · r/samsclub

Why do you guys sell whole pizzas made hours ago to customers?

"The very few times I've gotten a whole pie, it'll be stuff premade and left in the hot case for like an hour 30mins before it's in my hand. How can I tell? They put a sticker with the date and the pizza looks and taste hours old. Pizza is meant to be made to order and waited for."

ModernBERT: negative (conf 0.9999997) · aspects: customer service · product quality · store experience · trust 0.66

 [Open original on Reddit](#)

A · GPT-4o (Walmart Gateway)

Hi u/there, I'm sorry to hear about your experience with the pizzas. I understand how disappointing it must be to get something that feels less fresh than expected. While the hot case is meant to provide convenience, I agree that quality should never be sacrificed. If you'd like, feel free to DM me with the details of your store location so we can look into this and address it properly. – [Your Name]

B · Mistral 7B (local Ollama)

Hi u/there, I appreciate your concern about the quality of our prepared pizzas. It's important to us that our customers receive fresh and delicious food. The stickers with dates are there to help us manage the rotation of our products, ensuring they are consumed before they become stale. If you ever encounter a pizza that doesn't meet your expectations, please feel free to DM me the details of your order so I can look into it further. Thank you for taking the time to share your feedback.

C · Smart Composer (no-LLM)

Hi u/there — thanks for flagging this. completely understand the frustration around your order. Drop us a private message with the order # or store info and we'll take it from here. — Walmart Care

WHAT HAPPENS NEXT

One prompt fed to all three engines. Analyst picks the best draft, edits inline, clicks Save & Open Reddit → the posted reply is written to feedback and becomes the next post's top few-shot example.

Learning Loop — Feedback → Retraining

Every human action becomes a training signal — hot loop (live) + warm loop (monthly)

Label / aspect correction	Trust-score override	Posted reply	Lifecycle transition
Review & Validate	Review & Validate	Smart Reply Composer	Kanban board

feedback table — one row per human action, JSON payload, partition_key = analyst_id

HOT LOOP • in-context, live

- Every Generate Drafts click queries feedback
- SELECT ... WHERE kind = 'auto_reply_posted'
- ORDER BY created_at DESC LIMIT 5 → top 3 pairs
- Injected into the prompt as few-shot examples
- Effect visible on the next reply — no training

WARM LOOP • supervised, monthly

- Export corrections since last train date
- Append to Walmart-200 → Walmart-N (augmented)
- Rerun ModernBERT Stage-3 with 5-fold OOF CV
- New checkpoint wins only if OOF F1 improves
- Symlink flipped in config/models.yaml

Two failure modes eliminated by design

(1) hot loop uses in-context learning — no retrain, no catastrophic forgetting. (2) warm loop uses 5-fold OOF CV — no leakage between train and eval sets.

Post Explorer — Multi-Facet Search

Find the exact 20 posts an analyst needs to look at today

FACETS

- Sentiment · neg / neu / pos
- Confidence slider (0.50 – 1.00)
- Trust-score slider (0.00 – 1.00)
- Subreddit — multi-select from 25 tracked
- Aspect — 8-item retail taxonomy
- Date range — today / week / month / custom
- Full-text search — title + body

PER-POST CARD

Title + body excerpt · sentiment badge with confidence % · trust tier (H / M / L) · aspect chips · subreddit + timestamp · actions: Review · Add to Lifecycle · View Details.

CATALOG TODAY (live capture)

83 006 analysed posts · 25 subreddits · 8 aspects · results paginated at 50 / query.

Retail Sentiment INTELLIGENCE

Brand Health

Lifecycle

Insights

Review

Trust

Alerts

Posts

Pipeline

Notifications

Posts

Reddit · Dry-run

Post Explorer

Search the analyzed-post archive by subreddit, sentiment, aspect or trust score.

SUBREDDIT

SENTIMENT

ASPECT

RANGE

MIN TRUST

LIMIT

Search

e.g. walmart

All

e.g. pricing

All

0.0-1.0

50

Showing 50 of 83,602 matches (raise the Limit to see more)

negative

r/Frugal

↑ 1

8/1/2026, 11:15:45 PM

trust 0.54

Let go of nostalgia or loyalty for any brand or products, it is being used against you

The number one frugal tip I can think of is to stop being loyal or nostalgic for any brand or product. The moment a brand develops a reputation for being high quality is the moment that brand will begin exploiting that reputation by reducing the quality of their products and jacking up prices. That is just a natural consequence of a poor regulatory environment and a low-information consumer base. Exceptions to this are rare (e.g., maybe like LL Bean, Arizona Iced Tea, Ball Jars..., lol I can't think of...

View on Reddit

neutral

r/Frugal

↑ 1

8/1/2026, 9:55:53 PM

trust 0.56

How Would You Repair This Chair? Looking for a Quick and Budget-Friendly Solution.

View on Reddit

neutral

r/Frugal

↑ 1

8/1/2026, 8:21:19 PM

trust 0.54

Am I misremembering? Sams club membership

Did Sam's Club start with a group membership? You had to get a group (20? 40?) or join a group to get a membership? Do they still do that? I tried to join as the current special, but I'm not a new member since I had one 5 years ago :? Yes, it says 1 year, but ... its not. There's a senior member, you have to go thru [ID.me](a site i've never heard of. Can i walk in and join, show senior ID?

View on Reddit

negative

r/Frugal

↑ 1

8/1/2026, 7:56:57 PM

trust 0.54

Any frugal strategies for car rentals?

I'm flying out to a friend's wedding in a couple of months. Initially my plan was to use public transit for most of my travel while I'm there, which is still my primary option. But due to where they're having the wedding (a very rural area with zero public transport options) at some point I'm basically going to have to rent a car to get there. Sadly carpooling isn't really an option either due to a myriad of factors. So basically I'm gonna be stuck renting a car for at least a day or two, maybe three. Anyone...

View on Reddit

negative

r/Frugal

↑ 1

8/1/2026, 7:48:22 PM

trust 0.57

Which products are a good value at Whole Foods?

I just moved, and my primary grocery store is now going to be Aldi, with occasional trips to Trader Joe's for treats. However, both of those stores are a bit of a drive and don't open in time for me to stop in before work if I really need something. And then there is a Whole Foods right around the corner from me which opens at 7am. I normally avoid WF because the prices I have seen have all been outrageous, but I think I am going to have to utilize it at some point. My first day when I had no food, I...

View on Reddit

positive

r/Frugal

↑ 1

8/1/2026, 4:52:27 PM

trust 0.53

No Buy July 2026 - FINAL Check In - Week 4

A big thank you to everyone who participated! Hope you'll join in next July! How did the final week of No Buy July go for you? This is the thread to post how on how you did during the fourth week! Check in with your successes. Check in with your failures. Check in for accountability. Or simply follow along to see how others did during this final week! See what this challenge was all about by checking out the original post:

View on Reddit

neutral

r/Sparkdriver

↑ 1

8/1/2026, 2:53:16 PM

trust 0.53

No Service, No GPS at 10 PM

VISION CAPTION

Here's a description of the image: The image shows a muddy, overgrown dirt road leading away from a Walmart store. The road is heavily obscured by vegetation, suggesting a neglected or poorly maintained access path. There is no visible text or signage in the image.

View on Reddit

neutral

r/walmartop


↑ 1

8/1/2026, 2:52:43 PM

trust 0.53

Post Lifecycle — Kanban Workflow

Track a complaint from Triage to Resolved with SLA visibility



The screenshot shows the 'Post Lifecycle' interface within the Retail Sentiment Intelligence application. The interface is divided into three main columns: 'Ack & Reply Sent', 'Actionable Items', and 'Resolved'. Each column contains a list of posts with details such as sentiment (e.g., MEDIUM, LOW), source (e.g., reddit.com, amazon.com), and status (e.g., Action needed, Closed, Resolved). The 'Resolved' column also includes a 'Refresh' button. The interface is designed to provide a clear view of the complaint lifecycle and the actions taken to resolve them.

RESOLVE MODAL — 2-step flow

- Step 1 · save action note + optional LLM reply
- (a) Save & open Reddit → reply copied to clipboard
- (b) OR Resolve (no-reply) → close without posting
- Step 2 · paste on Reddit → Mark Resolved
- All transitions timestamped → SLA analytics

Insights & Competitor Analysis

Deterministic aggregation over the 8-aspect taxonomy — no LLM in the loop

Pain Points

Ranked table of competitor aspects by negative-ratio (top 10, floor ≥ 8 posts)

Walmart Comparison

Radar chart: Walmart neg-ratio vs competitors on the same aspects $\rightarrow$ gives delta

Recommendations

Template cards: priority tag + action angle chosen from ratio and delta thresholds

FORMULAS · what every number on the page is

1. Per-aspect negative rate (drives ranking + priority)

$$\text{negative_ratio}(\text{aspect}) = \text{neg} / \text{total} \quad (\text{floor: total} \geq 8 \text{ posts})$$

2. Walmart-vs-competitor gap (drives action angle)

$$\text{delta}(\text{aspect}) = \text{competitor negative_ratio} - \text{Walmart negative_ratio}$$

$\text{neg} / \text{total}$ = fraction of that aspect's competitor posts labelled negative by ModernBERT.

$\text{delta} > 0 \rightarrow$ competitors are worse than Walmart on this topic.

Zero LLM cost — every value is a Counter aggregate on top of ModernBERT + DeBERTa-NLI outputs.

THRESHOLDS · what the numbers become on the card

Priority tag (uses competitor negative_ratio)

- ratio $\geq 60\%$ $\rightarrow$ High priority
- ratio $\geq 40\%$ $\rightarrow$ Medium
- otherwise $\rightarrow$ Low

Action angle (uses delta)

- delta $> +5\%$ $\rightarrow$ 'Marketing angle' (Walmart is better)
- delta $< -5\%$ $\rightarrow$ 'Investigate' (Walmart is worse)
- within $\pm 5\%$ $\rightarrow$ 'Industry-wide friction'

Notification Centre

In-app mirror of the group-routed email / Slack digest

WHAT IT DOES

- Every P1 / P2 email or Slack alert is mirrored in the app
- Notifications grouped by subreddit ↔ business team
- Read / unread state persisted per analyst
- Deep-links open the post in Review & Validate
- Powered by the same alerts rows that drive email/Slack

PRIORITY THRESHOLDS

P1 — trust ≥ 0.70 and confidence ≥ 0.80
P2 — trust ≥ 0.50 and confidence ≥ 0.60
Source: config/models.yaml + dispatcher.py

GROUPS TODAY (live capture)

4 routing groups configured · DriverRelatedPost, GIF_, SAMSClub, W+ Membership · delivery status tracked per row in notification_log.

Brand Health

Lifecycle

Insights

Review

Trust

Alerts

Posts

Pipeline

Notifications

Internal analytics demo. Not affiliated with Walmart Inc. branding.

Notifications

Reddit · Dry-run

Notification Configuration

Configure which subreddit groups receive alerts for P1/P2 negative posts.

SENDER EMAIL

vishal.singh1@walmart.com

DriverRelatedPost

3 subreddits · 0 recipients · retailSentimentIntelligence

r/Sparkdriver

r/walmartgp

r/WalmartSparkDrivers

P1 + P2

GIF_

1 subreddit · 0 recipients · retailSentimentIntelligence

r/OGFBathroom

P1 + P2

SAMSClub

2 subreddits · 0 recipients · retailSentimentIntelligence

r/samsclub

r/Sams_Club

P1 + P2

W+ Membership

1 subreddit · 1 recipient · retailSentimentIntelligence

r/WalmartPlus

P1 + P2

Recent Notification Log

TIME	GROUP	CHANNEL	POST	STATUS
Aug 1, 03:28 PM	DriverRelatedPost	slack	reddit_1vccna3	FAILED
Aug 1, 03:28 PM	DriverRelatedPost	slack	reddit_1vcdntr	FAILED
Aug 1, 03:28 PM	DriverRelatedPost	slack	reddit_c_p8wfmvg	FAILED
Aug 1, 03:28 PM	DriverRelatedPost	slack	reddit_c_p8x6sz5	FAILED
Aug 1, 03:28 PM	DriverRelatedPost	slack	reddit_1vzb7h2	FAILED
Aug 1, 03:28 PM	DriverRelatedPost	slack	reddit_1vcapz3	FAILED
Aug 1, 03:28 PM	GIF_	slack	reddit_1vbvpms9	FAILED

ModernBERT — Final Training Results

3-stage curriculum, 5-fold out-of-fold CV, n = 200

Overall Macro-F1

0.6272 → 0.7642

Long posts (≥ 512 tok, n = 7)

5 / 7 → 7 / 7 correct

Uplift over baseline

+13.7 pts

3-stage curriculum


Stage	Data	Macro-F1	Δ vs baseline
Baseline (Twitter-RoBERTa)	off-the-shelf	0.6272	—
Stage 1	TweetEval sentiment (~45 k tweets)	0.6810	+5.4
Stage 2	GoEmotions-3class Reddit (~54 k, pseudo)	0.7285	+10.1
Stage 3 (final)	Walmart-200 hand-labelled	0.7642	+13.7

Per-length-bucket accuracy (5-fold OOF predictions)

Bucket (token count)	Baseline correct	ModernBERT correct	Recovered
Short-to-medium (< 512, n = 193)	138 / 193 (72 %)	159 / 193 (82 %)	+21
Long (≥ 512, n = 7, all negative)	5 / 7	7 / 7	+2

Vision Pipeline — Multi-Pass Payoff

Same model (Gemma 3 4B), same hardware, smarter calling strategy



Evaluation on 32 retail screenshots

Metric	Single-pass	Multi-pass	Change
Hallucination rate	50 %	0 %	eliminated
Text-extraction success	25 %	75 %	3× better
Retail-signal recall	40 %	81 %	2× better
Fabricated claims	many	0	eliminated
Latency / image (warm)	~5 s	~15 s	3× (accepted)

Trust-Score — End-to-End Evaluation

On the 200-post gold set, does the trust filter catch bad posts?

Low-trust share ($T < 0.4$)

15 %

Human-agreed (annotator cross-check)

12 / 15

Agreement rate

80 %

Weighted formula (decomposition shown inline in dashboard)

$$\text{trust_score} = 0.4 \times \text{metadata} + 0.3 \times \text{dedup} + 0.3 \times \text{llm_credibility}$$

DESIGN PRINCIPLE — flag, don't drop

Low-trust posts are surfaced with an explanatory chip (promotional / duplicate / thin-metadata) and can be overridden by an analyst in one click — the override is captured for the learning loop.

WHY IT'S DEFENSIBLE

- Every constant has a plain-English rationale in config/models.yaml
- Three named components → analyst can see WHY a post is low-trust
- 12 / 15 confirmed by human annotator on the cross-check
- LLM component only invoked in ambiguous zone (cost-controlled)

Evaluation Summary — Post-Midsem Numbers

All results reported against the same 200-post retail-Reddit gold set

Area	Metric	Value	Notes
Sentiment	Macro-F1 (final)	0.7642	+13.7 pts vs RoBERTa baseline
Sentiment	Long-bucket accuracy (≥ 512 tokens, n=7)	7 / 7 correct	Baseline 5 / 7 · all negative-class
Sentiment	Cross-validation	5-fold OOF	No test-set leakage
Vision	Hallucination rate	0 %	From 50 % on single-pass
Vision	Text extraction success	75 %	From 25 % on single-pass
Vision	Retail-signal recall	81 %	From 40 % on single-pass
Trust	Low-trust share (T < 0.4)	15 %	Of the 200-post set
Trust	Annotator agreement on low-trust	12 / 15	80 %
Ops	Subreddits tracked	25	Six community groups
Ops	Scheduler cadence	6 h + manual	asyncio lifespan task
Ops	Dashboard pages shipped	8	Brand Health · Alerts · Explorer · Review · Lifecycle · Insights · Pipeline · Notifications

Principal Contributions

What this dissertation delivered — end to end

C1

Offline-first RSI pipeline

Ingestion → trust → sentiment → aspects → vision → aggregation → alerts → dashboard. Zero API cost, modular, deployable to Azure.

C2

Fine-tuned ModernBERT (3-stage curriculum)

Macro-F1 0.6272 → 0.7642 overall; recovers all long posts (≥ 512 tok, $n = 7$): RoBERTa 5/7 vs ModernBERT 7/7 correct. 5-fold OOF cross-validation, offline training.

C3

Multi-pass vision technique on Gemma 3 4B

Hallucination 50 % → 0 %, extraction 25 % → 75 %. Adapts techniques from 5 recent VLM papers without vendor-blocked models.

C4

Interpretable trust score + admission gate

trust × confidence gate flags rather than drops low-credibility posts. Every constant traceable to English rationale in config.


C5

HITL learning-loop dashboard



Review & Validate, Lifecycle Kanban, Insights, Notification centre. Corrections & posted replies feed few-shot reply generation and future retraining.

Live Demo — Dashboard Walkthrough


One post → alert → review → reply → lifecycle → resolved




Brand Health




Alert Feed




Post Explorer



Review & Validate



Lifecycle Kanban



Insights

Conclusions — What Worked, What's Still Open

Honest read on 5 months of post-midsem work

WHAT WORKED

- 3-stage ModernBERT curriculum landed the sentiment win we set out for (Macro-F1 0.62 → 0.76 on 5-fold OOF).
- Long-post recovery came for free once we switched off truncation (≥ 1024 -tok context): 5/7 → 7/7 correct on the ≥ 512 -tok bucket.
- Removing the image on the final vision merge step ended the 50 % hallucination problem — the mechanism, not just the number, makes it credible.
- Trust score as flag-not-drop matched the human annotator on 12 / 15 low-trust posts. Analysts trust it because they can override it.
- HITL feedback table quietly turned into the most useful piece of infrastructure — drives few-shot on every click today, retraining tomorrow.

WHAT'S STILL OPEN (honest)

- 200-post gold set is small. Numbers are OOF-CV, but a bigger blind held-out set is needed before I'd claim generalisation.
- Vision eval is 32 images — enough to sanity-check the mechanism, not enough to claim production-grade quality.
- Long bucket has only 7 posts and they are all negative-class — 7/7 correct is evidence the truncation ceiling is gone, not proof of long-text mastery.
- Reply drafts still need an analyst edit — we track edit-distance but haven't shown it dropping consistently yet.
- Everything runs on one Mac. Multi-analyst concurrency and retraining automation aren't done — in the roadmap on next slide.

THE TAKEAWAY

Mechanism-level wins are what carry the argument; sample-size caveats are what the next horizon of work has to close.

Future Work — Track 1 • HITL Reduction (4 Horizons)

One of three tracks: Track 1 HITL • Track 2 Bilingual • Track 3 Multi-Platform (Tracks 2 & 3 on next slide)

Now → 1 month

HITL: 100 %

Grow the gold set

Grow the labelled gold set to ~500 posts using the HITL feedback the team is already producing; rerun 5-fold OOF to check the F1 delta.

1 → 3 months

HITL: 100 %

Automate the retrain loop

Wire the monthly ModernBERT retrain into a Cron / Airflow job with an OOF-F1 promotion gate — new checkpoints ship only when they beat the current one on held-out folds.

3 → 6 months

HITL: ~80 %

Walmart-trained model + shared DB

Fine-tune Mistral 7B on our own posted-reply corpus → a Walmart-trained replier that drops the paid GPT-4o dependency, runs fully in-house, and can be self-hosted. Move the feedback table to a shared managed database so multiple analysts can work concurrently.

6 → 12 months

HITL: ~30 % → eventually
0 %

Confidence-gated auto-reply

Introduce a confidence gate on the reply generator: top-confidence drafts are queued for auto-send and the analyst only reviews the mid- and low-confidence tail. As agreement stays high, gradually raise the auto-send threshold — HITL winds down to sampling / audit only.

GUIDING PRINCIPLE

HITL is the ladder, not the ceiling — every stage produces the training signal that lets the next stage automate more. Each horizon is admissible only when the previous one is measurably successful.

Future Work — Track 2 • Bilingual + Track 3 • Multi-Platform

The other two tracks that widen the pipeline's coverage while Track 1 winds HITL down

TRACK 2 • BILINGUAL → MULTILINGUAL

Why: Indian retail subreddits are ~30 % of Walmart-family mentions but the current langdetect → English-only filter silently drops them.

Concrete deliverables

- Extend langdetect to hi + en (later ta / te / bn); code-mixed 'Hinglish' handled with a mixed-script tokenizer.
- Ingestion-time translation via NLLB-200 (open-weights, 3.3 B) so downstream models stay English-only — or — fine-tune ModernBERT on native scripts.
- Extend aspect taxonomy for Indic retail idioms ("bill karo" for billing, "delivery kab" for late shipping).
- Native-speaker analyst validation loop — same HITL flow, new language pool.

Ships as part of Track 1 horizon 1 → 3 mo (bilingual retrain gate) and horizon 3 → 6 mo (Walmart-trained multilingual replier).

TRACK 3 • MULTI-PLATFORM (Reddit → X / Facebook / Instagram / app-stores)

Why: Reddit is deliberate + long-form; X is real-time viral; Instagram is image-heavy (vision pipeline shines); app-store reviews carry a star-rating ground-truth signal.

Concrete deliverables

- Ingestion adapters: X API v2, Meta Graph API (Facebook + Instagram), Google Play / App Store review scrapers — each behind a common RawPost interface.
- Per-platform pre-processing: X dedup by retweet lineage; Instagram routed through the vision pipeline first (caption + image); app-store reviews carry a native star_rating field.
- Extend priority classifier: X viral (retweets > 100) → auto P1; app-store 1-star surge → P1; Instagram brand-tagged post from > 10 k follower account → P1.
- Unified raw_posts.source_platform column (already stubbed) so Brand Health / Insights aggregate across platforms out of the box.

Risks: X API tiers (\$100 / mo min); Meta Graph needs page-level perms; Instagram image volume needs vision-batching.

Source Code & Deliverables

Repository

https://gecgithub01.walmart.com/v0s01jh/Retail_Sentiment_Intelligence

Deliverable	Path in repo	Contents
Final Report (PDF)	docs/Sem_4/final/FINAL_REPORT_VishalSingh_2020AA05641.pdf	10 chapters + appendices + refs
Final Report (LaTeX)	docs/Sem_4/final/latex/	Reproducible xelatex source
Pipeline core	src/pipeline.py, src/ingestion/, src/analysis/	6-layer async pipeline
Sentiment training	scripts/train_modernbert_sentiment.py	3-stage curriculum runner
Evaluation notebook	evaluation/trust_score_walmart200.ipynb	Reproducible 5-fold OOF eval
Dashboard	frontend/ (React + Vite + Tailwind)	8 pages, live via WebSocket
Reproduction	Appendix B of the report · start.sh	Clone → conda → start

REPRODUCE IN 3 STEPS

```
git clone <repo> → conda create -n rsi python=3.11 && pip install -r requirements.txt → ./start.sh
```

Thank You

Retail Sentiment Intelligence · Post Mid-Semester Presentation

Questions ?

[Repository](#)

[gecgithub01.walmart.com/v0s01jh/Retail_Sentiment_Intelligence](https://github.com/v0s01jh/Retail_Sentiment_Intelligence)

Report: docs/Sem_4/final/FINAL_REPORT_VishalSingh_2020AA05641.pdf

Vishal Singh · 2020AA05641 · BITS Pilani (WILP) · Walmart Global Tech, Bengaluru