

# Retail Sentiment Intelligence

Trust-Aware Sentiment Analysis of Retail Feedback using LLMs

BITS ZG628T • Dissertation • Post Mid-Semester Presentation

**Vishal Singh • ID No. 2020AA05641**

M.Tech (Artificial Intelligence & Machine Learning)

Birla Institute of Technology & Science, Pilani (WILP)

Dissertation work at Walmart Global Tech, Bengaluru

Supervisor: Mr. Varunendra Pratap Singh — Principal Software Engineer, Walmart Global Tech

Additional Examiner: Ms. Pradnya Kashikar — BITS Pilani

August 2026

# Agenda

Focus of today — post mid-semester work only

## 1 Since Mid-Sem — What's New

Post-midsem workstream map

## 2 Review & Validate (HITL)

Human-in-the-loop correction workflow

## 3 Smart Reply — Triple Draft

GPT-4o + Mistral 7B + Smart Composer

## 4 Smart Reply — Prompt & Few-Shot

Prompt template + past validated replies

## 5 Smart Reply — Worked Example

One complaint → three drafts

## 6 Learning Loop

Corrections + posted replies → retraining

## 7 Post Explorer

Multi-facet search across analysed posts

## 8 Post Lifecycle (Kanban)

Triage → Acknowledged → In-Progress → Resolved

## 9 Insights & Competitor Analysis

Strategic weekly summaries + cross-brand

## 10 Notification Centre

In-app + email + Slack, group-routed

## 11 ModernBERT — Final Results

Macro-F1 0.6272 → 0.7642 (+13.7 pts)

## 12 Vision — Multi-Pass Payoff

Hallucination 50% → 0%, extraction 25% → 75%

## 13 Trust-Score Evaluation

15% flagged; 12 of 15 confirmed by annotator

## 14 Storage — SQLite → Cosmos DB

Lift-and-shift ready

## 15 Live Demo

Dashboard walkthrough (screenshots)

## 16 Conclusions & Future Work

RO1-RO4 outcomes + roadmap

# Since Mid-Sem — What's New

Mid-sem shipped the pipeline; post-midsem made it usable for analysts

## Delivered at mid-sem

- 6-layer offline-first pipeline (25 subreddits)
- ModernBERT fine-tuning up to Stage 2 (F1 0.7285)
- Multi-pass vision with 0% hallucination on pilot
- Trust-score design + weighted formula
- Brand-Health + Alert Feed dashboard pages
- Group-routed email / Slack notifications

## Added post-midsem (this deck)

- Review & Validate — human-in-the-loop UI + backend
- Smart Reply Composer — dual-draft (rule + FLAN-T5)
- Learning-loop store — corrections & posted replies
- Post Explorer — multi-facet search
- Post Lifecycle — Kanban with 2-step Resolve
- Insights & Competitor Analysis page
- In-app Notification Centre mirroring the digest
- ModernBERT Stage 3 (final) — F1 0.7285 → 0.7642
- Trust-score end-to-end evaluation (n = 200)
- SQLite schema mirroring Cosmos DB partitioning


## Framing

Mid-sem answered can the pipeline produce trust-weighted sentiment? Post-midsem answers can an analyst act on it, correct it, and feed the corrections back into the model?

# Review & Validate — Human-in-the-Loop

Every prediction can be corrected; every correction is logged

- Queue sorted by priority (P1 first) · trust × confidence
- Analyst can override sentiment and aspect tags in one click
- Corrections written to feedback table with analyst + timestamp
- Generate Drafts → two reply options (rule + LLM)
- Analyst picks A or B, edits, and posts to Reddit
- Posted replies also captured → future few-shot pool



## Why it matters

Turns a passive dashboard into a training signal — corrections harvested here drive the ModernBERT re-training loop and reply few-shot pool.

# Smart Reply Composer — Triple Draft

One prompt, three generators; analyst picks the best draft

## Draft A

### GPT-4o

Walmart LLM Gateway

Strong reasoning, safest tone  
Guardrails via Walmart proxy  
~\$0.0002 / reply  
Fallback: direct OpenAI

## Draft B

### Mistral 7B-Instruct

Local Ollama (:11434)

Open-weights, free  
Runs fully offline  
~15 s warm latency  
Strong at retail jargon

## Draft C

### Smart Composer

Deterministic template

No LLM, always available  
Zero latency, zero cost  
Content-aware phrase pools  
Safety-net fallback

### Fallback logic — no draft is ever missing

- GPT-4o unavailable (no gateway key / consumer-ID / network) → Smart Composer draft in slot A
- Mistral unavailable (Ollama not running) → Smart Composer draft in slot B
- Smart Composer always produces slot C — guaranteed reply
- UI labels each card with the actual model used, plus an offline-fallback badge when applicable

# Smart Reply — Prompt Design & Few-Shot

Same prompt fed to GPT and Mistral; Smart Composer skips the LLM step

## Prompt template (built in `_build_reply_prompt`)

You are a senior Walmart customer-care analyst replying on Reddit.  
Write ONE reply to the customer below. Keep it 2-4 sentences,  
empathetic, specific to their complaint, no corporate jargon,  
no hashtags, no emojis. Do NOT promise refunds you can't verify;  
invite them to DM order details if action is needed. Sign off as  
a real person, not a brand.

Example customer post: {past\_post\_1}  
Example analyst reply: {past\_reply\_1}  
Example customer post: {past\_post\_2}  
Example analyst reply: {past\_reply\_2}

Subreddit: r/{subreddit}  
Customer ({author}) complaint about: {aspect\_1}, {aspect\_2}  
Customer post:  
{post\_title + post\_body, ≤ 1200 chars}

Reply:

### Few-shot source — feedback table

- Top-3 past validated replies for the same aspect
- Written by human analysts, already posted to Reddit
- Refreshed on every Generate click → adapts over time

### Style guardrails baked into the prompt

- 2 – 4 sentences, empathetic, specific
- No corporate jargon, no hashtags, no emojis
- No unverifiable refund promises → DM to continue
- Sign off as a person, not a brand

# Smart Reply — Worked Example (all three drafts real, live-captured)

raw\_posts.id = reddit\_1u2bgdw · r/samsclub · gateway ✓ · Ollama ✓ · Smart Composer ✓

## CUSTOMER COMPLAINT · r/samsclub

### Why do you guys sell whole pizzas made hours ago to customers?

"The very few times I've gotten a whole pie, it'll be stuff premade and left in the hot case for like an hour 30mins before it's in my hand. How can I tell? They put a sticker with the date and the pizza looks and taste hours old. Pizza is meant to be made to order and waited for."

ModernBERT: negative (conf 0.9999997) · aspects: customer service · product quality · store experience · trust 0.66

 [Open original on Reddit](#)

#### A · GPT-4o (Walmart Gateway)

Hi u/there, I'm sorry to hear about your experience with the pizzas. I understand how disappointing it must be to get something that feels less fresh than expected. While the hot case is meant to provide convenience, I agree that quality should never be sacrificed. If you'd like, feel free to DM me with the details of your store location so we can look into this and address it properly. – [Your Name]

#### B · Mistral 7B (local Ollama)

Hi u/there, I appreciate your concern about the quality of our prepared pizzas. It's important to us that our customers receive fresh and delicious food. The stickers with dates are there to help us manage the rotation of our products, ensuring they are consumed before they become stale. If you ever encounter a pizza that doesn't meet your expectations, please feel free to DM me the details of your order so I can look into it further. Thank you for taking the time to share your feedback.

#### C · Smart Composer (no-LLM)

Hi u/there — thanks for flagging this. completely understand the frustration around your order. Drop us a private message with the order # or store info and we'll take it from here. — Walmart Care

One prompt fed to all three engines. Analyst picks the best draft, edits inline, clicks Save & Open Reddit → the posted reply is written to feedback and becomes the next post's top few-shot example.

# Learning Loop — Feedback → Retraining

Every human action becomes a training signal — hot loop (live) + warm loop (monthly)

| Label / aspect correction | Trust-score override | Posted reply         | Lifecycle transition |
|---------------------------|----------------------|----------------------|----------------------|
| Review & Validate         | Review & Validate    | Smart Reply Composer | Kanban board         |

feedback table (SQLite / Cosmos DB) — one row per human action, JSON payload, partition\_key = analyst\_id

## HOT LOOP • in-context, live

- Every Generate Drafts click queries feedback
- SELECT ... WHERE kind = 'auto\_reply\_posted'
- ORDER BY created\_at DESC LIMIT 5 → top 3 pairs
- Injected into the prompt as few-shot examples
- Effect visible on the next reply — no training

## WARM LOOP • supervised, monthly

- Export corrections since last train date
- Append to Walmart-200 → Walmart-N (augmented)
- Rerun ModernBERT Stage-3 with 5-fold OOF CV
- New checkpoint wins only if OOF F1 improves
- Symlink flipped in config/models.yaml

## Two failure modes eliminated by design

(1) hot loop uses in-context learning — no retrain, no catastrophic forgetting. (2) warm loop uses 5-fold OOF CV — no leakage between train and eval sets.



# Post Explorer — Multi-Facet Search

Find the exact 20 posts an analyst needs to look at today

## Facets

- Sentiment · neg / neu / pos
- Confidence slider (0.50 – 1.00)
- Trust-score slider (0.00 – 1.00)
- Subreddit — multi-select from 25 tracked
- Aspect — 8-item retail taxonomy
- Date range — today / week / month / custom
- Full-text search — title + body



Brand Health

Lifecycle

Insights

Review

Trust

Alerts

**Posts**

Pipeline

Notifications

Posts

Reddit · Dry-run

Post Explorer

Search the analyzed-post archive by subreddit, sentiment, aspect or trust score.

SUBREDDIT

SENTIMENT

ASPECT

RANGE

MIN TRUST

LIMIT

Search

e.g. walmart

All

e.g. pricing

All

0.0-1.0

50

Showing 50 of 83,006 matches (raise the Limit to see more)

neutral

r/WalmartEmployees

↑ 1

7/31/2026, 3:33:40 PM

trust 0.55

onepay

newbie here. how does instapay work on the onepay app?

View on Reddit

negative

r/AmazonPrime

↑ 1

7/31/2026, 3:33:25 PM

trust 0.66

Amazon promised me a \$150 shipping refund in chat to make me return an item, then refused to pay. Can they just lie like that?

Hey everyone, I need some advice or to know if anyone else has dealt with this nightmare. I recently had to return an item to Amazon internationally. When I was setting up the return through live chat, the support agent explicitly told me: "You can use any carrier to return the item. You need to pay to return the item. But the good thing is, we can help you with the refund for return shipping charges up to \$150." Because of that exact, written promise, I went ahead, paid out of pocket, and shipped the...

View on Reddit

negative

r/eCommerce

↑ 1

7/31/2026, 3:27:20 PM

trust 0.66

Struggling with cancelled parcels in fulfillment

We've been seeing a rise in cancelled orders on TikTok, Lazada, and Shopee that are already processed and packed in our fulfillment area. The problem is, by the time we realize they're cancelled, the items are fully packed— meaning wasted packaging materials and extra work. I'm trying to figure out how to catch these cancellations earlier in the workflow so we can save on cartons, fillers, and waybills. Has anyone here dealt with this? How do you identify cancelled items before they even reach...

View on Reddit

positive

r/CustomerService

↑ 1

7/31/2026, 3:24:38 PM

trust 0.66

Dear ppl who call in to customer service


I am a call center agent, I actually enjoy my job! but for those who call into anywhere yelling and fussing about goofy shyt. or a package coming a day late, or ur broke ass begging for a discount... get it together 🤔🤔 I hope you get hung up on until you know how to talk to people. I hang up on customers all the time🤔 soon as

## Per-post card

Title + body excerpt · sentiment badge with confidence % · trust tier (H / M / L) · aspect chips · subreddit + timestamp · actions: Review · Add to Lifecycle · View Details.

# Post Lifecycle — Kanban Workflow

Track a complaint from Triage to Resolved with SLA visibility



The screenshot shows the 'Post Lifecycle' dashboard in the Retail Sentiment Intelligence application. The dashboard is divided into three main columns: 'Ack & Reply Sent', 'Actionable Items', and 'Resolved'. Each column contains a list of posts with their status, priority, and source. The 'Ack & Reply Sent' column has 2 items, 'Actionable Items' has 3, and 'Resolved' has 12. The dashboard also includes a 'Refresh' button and a 'Reddit - Dry-run' indicator.

**Retail Sentiment Intelligence**

**Post Lifecycle**

Auto-created from confidently-negative posts. Slack and email notifications fire on entry (dry-run by default).

**Ack & Reply Sent** (2 items)

- MEDIUM (untitled) 23h ago, r/OGPBackroom, customer service. Action needed, Close.
- MEDIUM Can anyone explain the ATC role to me? 23h ago, r/OGPBackroom, store experience. Action needed, Close.

**Actionable Items** (3 items)

- MEDIUM I think I fucked up today 1d ago, r/OGPBackroom, customer service. Resolve.
- MEDIUM Seeking Advice: Walmart Career Growth and Promotion Challenges After 2 Years 9d ago, r/WalmartCanada, workforce\_hr. Resolve.
- MEDIUM (untitled) 18d ago, r/WalmartSparkDrivers, online/app. Resolve.

**Resolved** (12 items)

- MEDIUM (untitled) 23h ago, r/WalmartSparkDrivers, online/app.
- LOW (untitled) 23h ago, r/amazonprime.
- LOW Rated 1 star a passibe aggressive agent who was unhelpful and my amazon prime got... 23h ago, r/AmazonPrime.
- LOW Dear Corporate 1d ago, r/walmartgp.
- MEDIUM How do I convince associates to sprint out the door 1d ago, r/walmartgp, customer service.

## Resolve modal — 2-step flow

- Step 1 · save action note + optional LLM reply
- (a) Save & open Reddit → reply copied to clipboard
- (b) OR Resolve (no-reply) → close without posting
- Step 2 · paste on Reddit → Mark Resolved
- All transitions timestamped → SLA analytics

# Insights & Competitor Analysis

Strategic view — issue rankings, cross-brand pulse, LLM summaries

## Priority-Negatives

Top issues ranked by  
volume × severity × recency  
by aspect

## Competitor Pulse

Walmart vs Costco,  
Target, Amazon on shared aspects

## Weekly LLM Summaries

Natural-language weekly  
digest + action items  
+ emerging topics

## Aspect Drilldown

8-aspect taxonomy  
per-aspect sentiment trend  
+ representative posts

Screenshot on Slide 17 (Live Demo grid) — insights\_competitor page.

# Notification Centre

In-app mirror of the group-routed email / Slack digest

- Every P1 / P2 email or Slack alert is mirrored in the app
- Notifications grouped by subreddit ↔ business team
- Read / unread state persisted per analyst
- Deep-links open the post directly in Review & Validate
- Powered by the same alerts rows that drive email/Slack



Brand Health

Lifecycle

Insights

Review

Trust

Alerts

Posts

Pipeline

Notifications

Notifications

Reddit - Dry-run

Notification Configuration

Configure which subreddit groups receive alerts for P1/P2 negative posts.

SENDER EMAIL

vishal.singh1@walmart.com

DriverRelatedPost

3 subreddits · 0 recipients · retailSentimentIntelligence

r/Sparkdriver

r/walmartgo

r/WalmartSparkDrivers

P1 + P2

GIF\_

1 subreddit · 0 recipients · retailSentimentIntelligence

r/OGPBackroom

P1 + P2

SAMSClub

2 subreddits · 0 recipients · retailSentimentIntelligence

r/samsclub

r/Sams\_Club

P1 + P2

W+ Membership

1 subreddit · 1 recipient · retailSentimentIntelligence

r/WalmartPlus

P1 + P2

Recent Notification Log


| TIME             | GROUP             | CHANNEL | POST             | STATUS |
|------------------|-------------------|---------|------------------|--------|
| Jul 31, 04:15 PM | DriverRelatedPost | slack   | reddit_1vb5jl7   | FAILED |
| Jul 31, 04:15 PM | DriverRelatedPost | slack   | reddit_1vblyrt   | FAILED |
| Jul 31, 04:15 PM | DriverRelatedPost | slack   | reddit_c_p8qqv16 | FAILED |

### Priority thresholds

P1 — trust ≥ 0.70 and confidence ≥ 0.80  
P2 — trust ≥ 0.50 and confidence ≥ 0.60  
Source: config/models.yaml + dispatcher.py

# ModernBERT — Final Training Results

3-stage curriculum, 5-fold out-of-fold CV, n = 200



## 3-stage curriculum


| Stage                      | Data                                     | Macro-F1 | Δ vs baseline |
|----------------------------|------------------------------------------|----------|---------------|
| Baseline (Twitter-RoBERTa) | off-the-shelf                            | 0.6272   | —             |
| Stage 1                    | TweetEval sentiment (~45 k tweets)       | 0.6810   | +5.4          |
| Stage 2                    | GoEmotions-3class Reddit (~54 k, pseudo) | 0.7285   | +10.1         |
| Stage 3 (final)            | Walmart-200 hand-labelled                | 0.7642   | +13.7         |

## Per-length-bucket F1

| Bucket                   | Baseline | ModernBERT | Δ     |
|--------------------------|----------|------------|-------|
| Short (< 64 tokens)      | 0.7500   | 0.7800     | +3.0  |
| Medium (64 - 256 tokens) | 0.6500   | 0.7400     | +9.0  |
| Long (> 256 tokens)      | 0.2778   | 1.0000     | +72.2 |

# Vision Pipeline — Multi-Pass Payoff

Same model (Gemma 3 4B), same hardware, smarter calling strategy



## Evaluation on 32 retail screenshots

| Metric                  | Single-pass | Multi-pass | Change        |
|-------------------------|-------------|------------|---------------|
| Hallucination rate      | 50 %        | 0 %        | eliminated    |
| Text-extraction success | 25 %        | 75 %       | 3× better     |
| Retail-signal recall    | 40 %        | 81 %       | 2× better     |
| Fabricated claims       | many        | 0          | eliminated    |
| Latency / image (warm)  | ~5 s        | ~15 s      | 3× (accepted) |

# Trust-Score — End-to-End Evaluation

On the 200-post gold set, does the trust filter catch bad posts?

Low-trust share ( $T < 0.4$ )

15 %

Human-agreed (annotator cross-check)

12 / 15

Agreement rate

80 %

Weighted formula (decomposition shown inline in dashboard)

$$\text{trust\_score} = 0.4 \times \text{metadata} + 0.3 \times \text{dedup} + 0.3 \times \text{llm\_credibility}$$

## Design principle — flag, don't drop

Low-trust posts are surfaced with an explanatory chip ("promotional / duplicate / thin-metadata") and can be overridden by an analyst in one click — the override is captured for the learning loop. Every constant in the formula has a stakeholder-arguable English rationale in config/models.yaml.

# Storage — SQLite → Cosmos DB Lift-and-Shift

Same schema, same partition keys, no calling-code change

| Container        | Partition Key | Key Fields                                     | Purpose                    |
|------------------|---------------|------------------------------------------------|----------------------------|
| raw_posts        | /subreddit    | id, title, body, author_hash, created_utc      | Ingested (privacy-safe)    |
| analyses         | /subreddit    | post_id, sentiment, confidence, aspects, trust | AI analysis results        |
| aggregates       | /time_window  | subreddit, window, metrics_json                | Pre-computed KPIs          |
| alerts           | /severity     | type, aspect, threshold_breached               | Triggered anomalies        |
| feedback         | /analyst_id   | post_id, correction, reply_text                | HITL corrections + replies |
| notification_log | /group_id     | post_id, channel, status, sent_at              | Delivery audit trail       |

## Design choices

- SQLite in WAL mode for hot dev data
- data JSON column mirrors Cosmos doc
- Pluggable StorageBackend interface
- Nightly backup + JSONL cost ledger

## Migration story

- Swap SQLiteBackend → CosmosBackend
- Same CREATE TABLE mapped to container spec
- Partition keys already match production
- Zero change to pipeline / dashboard code



# Evaluation Summary — Post-Midsem Numbers

All results reported against the same 200-post retail-Reddit gold set

| Area      | Metric                           | Value        | Notes                                                                                       |
|-----------|----------------------------------|--------------|---------------------------------------------------------------------------------------------|
| Sentiment | Macro-F1 (final)                 | 0.7642       | +13.7 pts vs RoBERTa baseline                                                               |
| Sentiment | Long-post F1 (> 256 tokens)      | 1.0000       | From 0.2778 baseline                                                                        |
| Sentiment | Cross-validation                 | 5-fold OOF   | No test-set leakage                                                                         |
| Vision    | Hallucination rate               | 0 %          | From 50 % on single-pass                                                                    |
| Vision    | Text extraction success          | 75 %         | From 25 % on single-pass                                                                    |
| Vision    | Retail-signal recall             | 81 %         | From 40 % on single-pass                                                                    |
| Trust     | Low-trust share (T < 0.4)        | 15 %         | Of the 200-post set                                                                         |
| Trust     | Annotator agreement on low-trust | 12 / 15      | 80 %                                                                                        |
| Ops       | Subreddits tracked               | 25           | Six community groups                                                                        |
| Ops       | Scheduler cadence                | 6 h + manual | asyncio lifespan task                                                                       |
| Ops       | Dashboard pages shipped          | 8            | Brand Health · Alerts · Explorer · Review · Lifecycle · Insights · Pipeline · Notifications |

# Principal Contributions

What this dissertation delivered — end to end

C1

## Offline-first RSI pipeline

Ingestion → trust → sentiment → aspects → vision → aggregation → alerts → dashboard. Zero API cost, modular, deployable to Azure.

C2

## Fine-tuned ModernBERT (3-stage curriculum)

Macro-F1 0.6272 → 0.7642 overall; 0.28 → 1.00 on long posts. 5-fold OOF cross-validation, offline training.

C3

## Multi-pass vision technique on Gemma 3 4B

Hallucination 50 % → 0 %, extraction 25 % → 75 %. Adapts techniques from 5 recent VLM papers without vendor-blocked models.

C4

## Interpretable trust score + admission gate

trust × confidence gate flags rather than drops low-credibility posts. Every constant traceable to English rationale in config.


C5

## HITL learning-loop dashboard


Review & Validate, Lifecycle Kanban, Insights, Notification centre. Corrections & posted replies feed few-shot reply generation and future retraining.

# Live Demo — Dashboard Walkthrough


One post → alert → review → reply → lifecycle → resolved




## Brand Health




## Review & Validate




## Alert Feed



## Lifecycle Kanban



## Post Explorer



## Insights

# Conclusions — Research Questions Answered

**RQ1**

**Can a fine-tuned encoder beat baselines on Reddit retail sentiment?**

YES — ModernBERT 3-stage curriculum raised Macro-F1 from 0.6272 → 0.7642 overall and from 0.2778 → 1.0000 on long posts.

**RQ2**

**Can a compliant open-weights VLM extract structured retail signal from screenshots?**

YES — Multi-pass pipeline on Gemma 3 4B reduced hallucination from 50 % → 0 % and lifted text extraction from 25 % → 75 %.

**RQ3**

**Can a defensible trust score filter low-credibility posts without silent drops?**

YES — Interpretable 3-part score flagged 15 % of posts; 12 of 15 confirmed by human annotator. All low-trust posts remain visible and analyst-overridable.

**RQ4**

**Can a HITL workflow produce a re-training signal from analyst corrections?**

YES — Every correction and posted reply is logged in the `feedback` table → feeds ModernBERT Stage-3 augmentation and FLAN-T5 few-shot pool.

# Recommendations — Immediate Follow-Ons

## Monthly retraining cadence

Use the HITL feedback store as the incremental labelling stream and re-run the 3-stage curriculum monthly.

## Bilingual pass (Hindi + English)

Validate the retail taxonomy with native-speaker analysts before extending to Indian retail communities.

## Per-team SLA dashboards

Extend the Post Lifecycle SLA analytics into per-team dashboards once analyst volume passes ~50 posts / day.

# Future Work — Model + Product

Directions grouped by layer of the stack

## Model-level

- Joint sentiment + aspect head on a shared encoder (~30 % inference saving)
- Distil ModernBERT to ~50 M-parameter student → CPU-only edge inference
- Reasoning-augmented VLM captioning (LLaVA-Next 1.6 B or SmolVLM)
- 3-seed ensemble for tighter F1 variance
- Blind 25-post recheck for defensibility

## Product-level

- Auto-reply confidence gate — promote LLM drafts to "queued to send" when analyst edit-distance < threshold
- Predictive P1 forecast — seasonal decomposition on daily counts (24–48 h ahead)
- Bilingual (Hindi + English) taxonomy
- Slack-bot inline notification responses
- Automated retraining pipeline hook

# Future Work — Operational

Path from local prototype to Walmart-production service

## Kubernetes CronJob

Deploy the pipeline as a scheduled Kubernetes CronJob backed by a managed Postgres instance for multi-analyst concurrency.

## Azure Cosmos DB migration

Swap SQLiteBackend for CosmosBackend (schema already mirrors partition design) — no calling-code change.

## Walmart ticketing integration

P1 alerts open cases directly in Walmart's internal ticketing so the analyst never leaves the workflow.

## Broader ingestion

Extend beyond Reddit to Twitter / X, YouTube comments, and app-store reviews on the same trust + sentiment stack.

# Source Code & Deliverables

## Repository

[https://gecgithub01.walmart.com/v0s01jh/Retail\\_Sentiment\\_Intelligence](https://gecgithub01.walmart.com/v0s01jh/Retail_Sentiment_Intelligence)

| Deliverable          | Path in repo                                              | Contents                        |
|----------------------|-----------------------------------------------------------|---------------------------------|
| Final Report (PDF)   | docs/Sem_4/final/FINAL_REPORT_VishalSingh_2020AA05641.pdf | 10 chapters + appendices + refs |
| Final Report (LaTeX) | docs/Sem_4/final/latex/                                   | Reproducible xelatex source     |
| Pipeline core        | src/pipeline.py, src/ingestion/, src/analysis/            | 6-layer async pipeline          |
| Sentiment training   | scripts/train_modernbert_sentiment.py                     | 3-stage curriculum runner       |
| Evaluation notebook  | evaluation/trust_score_walmart200.ipynb                   | Reproducible 5-fold OOF eval    |
| Dashboard            | frontend/ (React + Vite + Tailwind)                       | 8 pages, live via WebSocket     |
| Reproduction         | Appendix B of the report · start.sh                       | Clone → conda → start           |



# Thank You

Retail Sentiment Intelligence · Post Mid-Semester Presentation

## Questions ?

[Repository](#)

[gecgithub01.walmart.com/v0s01jh/Retail\\_Sentiment\\_Intelligence](https://github.com/v0s01jh/Retail_Sentiment_Intelligence)

Report: docs/Sem\_4/final/FINAL\_REPORT\_VishalSingh\_2020AA05641.pdf

Vishal Singh · 2020AA05641 · BITS Pilani (WILP) · Walmart Global Tech, Bengaluru